

# **Sparser, better, faster, stronger**

Automatic differentiation with a lot of zeros

---

Guillaume Dalle\* – LVMT, École des Ponts ([gdalle.github.io](https://gdalle.github.io))

Laboratoire Jean Kuntzmann, 19.06.2025

---

\*joint work with Adrian Hill, Alexis Montoison and Assefaw Gebremedhin

# Agenda

1. Motivation
2. Automatic differentiation
3. Leveraging sparsity
4. Implementation

# Motivation



# Newton's method

## Root-finding

Solve  $F(x) = 0$  by iterating

$$x_{t+1} = x_t - \underbrace{[\partial F(x_t)]^{-1}}_{\text{Jacobian}} F(x_t)$$

## Optimization

Solve  $\min_x f(x)$  by iterating

$$x_{t+1} = x_t - \underbrace{[\nabla^2 f(x_t)]^{-1}}_{\text{Hessian}} \nabla f(x_t)$$

**Linear system** involving a derivative matrix  $A$ .

# Implicit differentiation

Differentiate  $x \mapsto y(x)$  knowing **conditions**  $c(x, y(x)) = 0$ .

Applications: fixed-point iterations, optimization problems.

Implicit function theorem:

$$\frac{\partial}{\partial x} c(x, y(x)) + \frac{\partial}{\partial y} c(x, y(x)) \cdot \partial y(x) = 0$$

$$\partial y(x) = - \underbrace{\left[ \frac{\partial}{\partial y} c(x, y(x)) \right]^{-1}}_{\text{Jacobian}} \frac{\partial}{\partial x} c(x, y(x))$$

**Linear system** involving a derivative matrix  $A$ .

# Linear systems of equations

How to solve  $Au = v$ ?

**Direct method** (LU, Cholesky)

1. Decompose the matrix  $A$ .
2. Get an exact solution by substitution.

Requires storing  $A$  explicitly.

**Iterative method** (CG, GMRES)

1. Rephrase as  $\min_u \|Au - v\|^2$ .
2. Get an approximate solution.

Only requires matrix-vector products  $u \mapsto Au$ .

# Conventional wisdom

- Jacobian and Hessian matrices are **too large** to compute or store
- We can only access linear maps  $u \mapsto Au$  (JVPs, VJPs, HVPs)
- Linear systems  $A^{-1}v$  must be solved with **iterative methods**
- Downsides: each iteration is expensive, convergence is tricky

# The benefits of sparsity

- Jacobian and Hessian matrices have **mostly zero coefficients**
- We can compute and store  $A$  explicitly
- Linear systems  $A^{-1}v$  can be solved with iterative **or direct** methods
- Upsides: faster iterations or exact solves, efficient linear algebra



# Automatic differentiation

---

# Differentiation

Given  $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$ , its **differential**  $\partial f(x)$  is the **linear map** that best approximates  $f$  around  $x$ :

$$f(x + u) = f(x) + \partial f(x)[u] + o(u)$$

It can be represented by the Jacobian **matrix**, which I will denote by  $\partial_{\text{mat}} f(x)$  instead.

# Numeric differentiation

input	output
program computing the function $x \mapsto f(x)$	approximation of the differential with the same program $\partial f(x)[u] \approx \frac{f(x + \varepsilon u) - f(x)}{\varepsilon}$

# Automatic / algorithmic differentiation

input	output
program computing the function $x \mapsto f(x)$	new program computing the exact differential $(x, u) \mapsto \partial f(x)[u] \text{ or } \partial f(x)^*[u]$

# AD under the hood

Two ingredients only:

1. hardcode basic derivatives (+, ×, exp, log, ...)
2. handle composition  $f = g \circ h$

# Composition

For a function  $f = g \circ h$ , the **chain rule** gives its differential:

$$\text{standard} \quad \partial f(x) = \partial g(h(x)) \circ \partial h(x)$$

$$\text{adjoint} \quad \partial f(x)^* = \partial h(x)^* \circ \partial g(h(x))^*$$

These linear maps apply as follows:

$$\text{forward} \quad \partial f(x) : U \xrightarrow{\partial h(x)} V \xrightarrow{\partial g(h(x))} W$$

$$\text{reverse} \quad \partial f(x)^* : U \xleftarrow{\partial h(x)^*} V \xleftarrow{\partial g(h(x))^*} W$$

# Why linear maps?

The chain rule has a matrix equivalent:

$$\partial_{\text{mat}}(g \circ h)(x) = \partial_{\text{mat}}g(h(x)) \cdot \partial_{\text{mat}}h(x)$$

$$\partial_{\text{mat}}(g \circ h)(x)^T = \partial_{\text{mat}}h(x)^T \cdot \partial_{\text{mat}}g(h(x))^T$$

Working with linear maps avoids allocation and manipulation of **intermediate Jacobian matrices**.

Essential for neural networks!

# Pocket AD

```
# Basic rules
```

```
using LinearAlgebra
```

```
A, b = rand(2, 3), rand(2)
```

```
residuals(x) = A * x - b
```

```
∂(::typeof(residuals)) = x → (u → A * u) #  $\mathbb{R}^3 \rightarrow \mathbb{R}^2$ 
```

```
∂T(::typeof(residuals)) = x → (v → adjoint(A) * v) #  $\mathbb{R}^2 \rightarrow \mathbb{R}^3$ 
```

```
sqnorm(r) = sum(abs2, r)
```

```
∂(::typeof(sqnorm)) = r → (v → dot(2r, v)) #  $\mathbb{R}^2 \rightarrow \mathbb{R}$ 
```

```
∂T(::typeof(sqnorm)) = r → (w → 2r .* w) #  $\mathbb{R} \rightarrow \mathbb{R}^2$ 
```



# Pocket AD

# Composition

```
function ∂(f :: ComposedFunction)
    g, h = f.outer, f.inner
    return x → ∂(g)(h(x)) ∘ ∂(h)(x)
end
```

```
function ∂⊤(f :: ComposedFunction)
    g, h = f.outer, f.inner
    return x → ∂⊤(h)(x) ∘ ∂⊤(g)(h(x))
end
```

# Pocket AD

```
julia> import ForwardDiff as FD, Zygote
```

```
julia> f = sqnorm ∘ residuals;
```

```
julia> x, Δx = rand(3), [1, 0, 0];
```

```
julia> ∂(f)(x)(Δx) # partial derivative  
0.8691056836969242
```

```
julia> ∂T(f)(x)(1) # gradient  
3-element Vector{Float64}:  
 0.8691056836969242  
 0.9973491983376236  
 0.5768822265195823
```

```
julia> FD.derivative(t → f(x + t * Δx), 0)  
0.8691056836969242
```

```
julia> Zygote.gradient(f, x)[1]  
3-element Vector{Float64}:  
 0.8691056836969242  
 0.9973491983376236  
 0.5768822265195823
```

# Two modes

Forward-mode AD computes Jacobian-Vector Products (**JVPs**) = “pushforward” of an **input perturbation**:

$$u \mapsto \partial f(x)[u] = Ju$$

Reverse-mode AD computes Vector-Jacobian Products (**VJPs**) = “pullback” of an **output sensitivity**:

$$v \mapsto \partial f(x)^*[v] = J^T v = v^T J$$

**Theorem (Baur-Strassen): cost of 1 JVP  
or VJP  $\propto$  cost of 1 function evaluation**

# What about gradients?

Reverse mode computes **gradients for roughly the same cost** as the function itself:

$$\nabla f(x) = \partial f(x)^*[1]$$

Makes deep learning possible.

The devil is in the details: higher memory footprint.




Figure 1: Blondel & Roulet (2024)

# What about second order?

The Hessian matrix is the Jacobian matrix of the gradient function.

A Hessian-Vector Product (HVP) can be computed as the JVP of a VJP, in **forward-over-reverse mode**<sup>2</sup>:

$$\nabla^2 f(x)[v] = \partial(\nabla f)(x)[v] = \partial(\partial^* f(x)[1])[v]$$

---

<sup>2</sup>Pearlmutter (1994)

# Leveraging sparsity

---

# From maps to matrices

To compute the Jacobian matrix  $J$  of a composition  $f : \mathbb{R}^m \longrightarrow \mathbb{R}^n$ :

- ~~product of intermediate Jacobian matrices~~
- reconstruction from several JVPs or VJPs

	<b>forward mode</b>	<b>reverse mode</b>
idea	1 JVP gives 1 column	1 VJP gives 1 row
formula	$J_{\cdot,j} = \partial f(x)[e_j]$	$J_{i,\cdot} = \partial f(x)^*[e_i]$
cost	$n$ JVPs (input dimension)	$m$ VJPs (output dimension)



# Using fewer products

When the Jacobian is sparse, we can compute it faster<sup>3</sup>.

If columns  $j_1, \dots, j_k$  of  $J$  are structurally **orthogonal** (their nonzeros never overlap), we deduce them all from a single JVP:

$$J_{j_1} + \dots + J_{j_k} = \partial f(x)[e_{j_1} + \dots + e_{j_k}]$$

Once we have grouped columns, sparse AD has two steps:

3. one JVP for each group  $c = \{j_1, \dots, j_k\}$
4. decompression into individual columns  $j_1, \dots, j_k$

---

<sup>3</sup>Curtis et al. (1974)

# Two preliminary steps

When grouping columns, we want to


- guarantee orthogonality (correctness)  $\Rightarrow$  pattern detection
- form the smallest number of groups (efficiency)  $\Rightarrow$  coloring

preparation	execution
1. pattern detection 2. coloring	3. matrix-vector products 4. decompression


The preparation phase can be **amortized** across several inputs.

# The gist in one slide

(a) AD code transformation



(b) Standard AD Jacobian computation



(c) ASD Jacobian computation

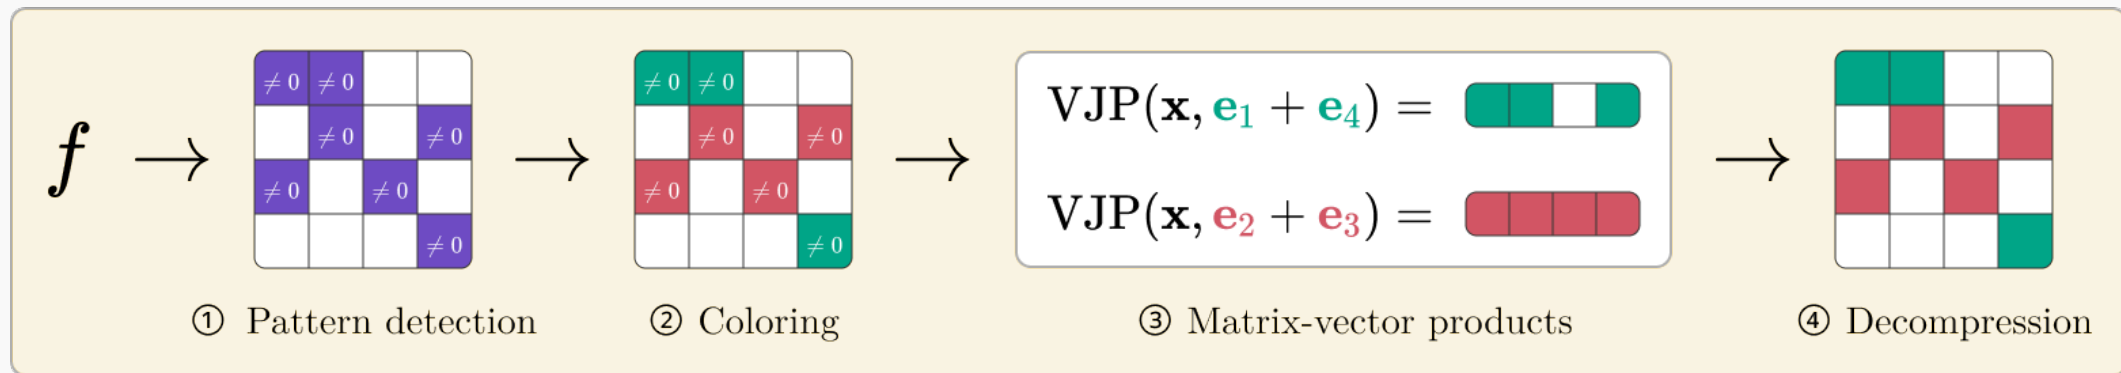
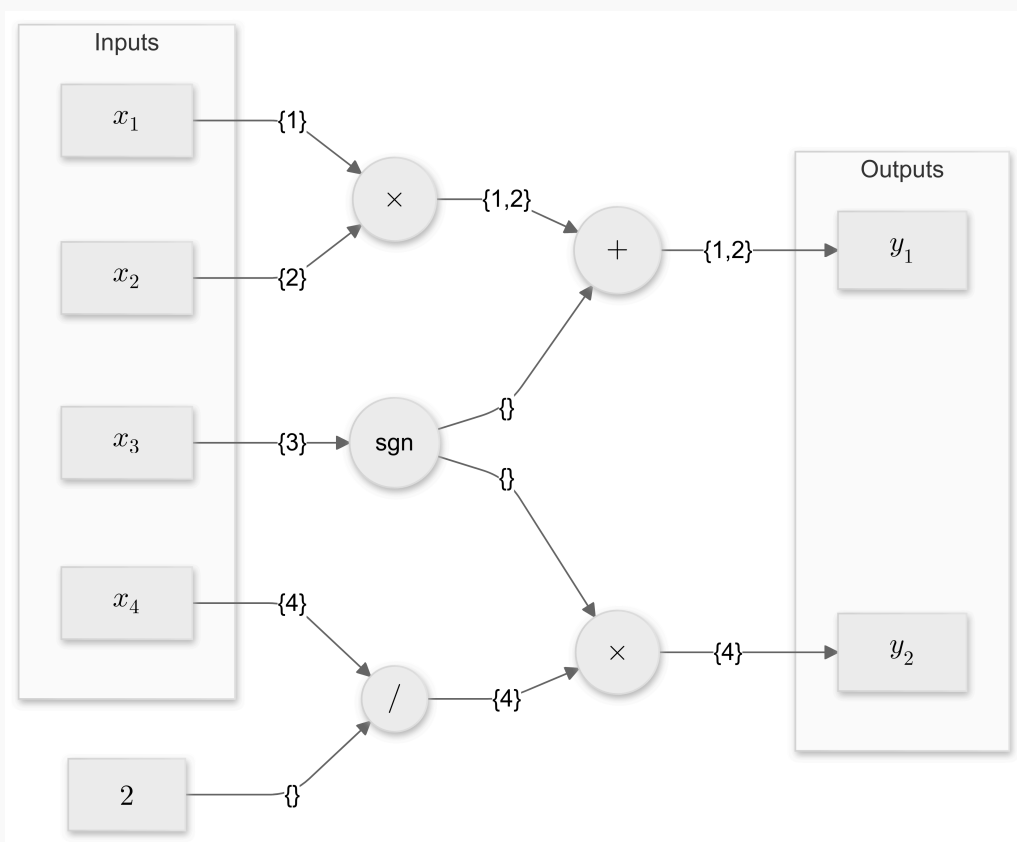


Figure 2: Hill & Dalle (2025)

# Tracing dependencies in the computation graph



Computation graph for

$$y_1 = x_1 x_2 + \text{sign}(x_3)$$

$$y_2 = \text{sign}(x_3) \times \left( \frac{x_4}{2} \right)$$

Its Jacobian matrix will have 3 nonzero coefficients:

$$\begin{pmatrix} 1 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 \end{pmatrix}$$

# Pocket pattern detection

```
import Base: +, *, /, sign
```

```
struct Tracer  
    indices :: Set{Int}  
end
```

```
Tracer() = Tracer(Set{Int}())
```

```
+(a :: Tracer, b :: Tracer) = Tracer(a.indices ∪ b.indices)
```

```
*(a :: Tracer, b :: Tracer) = Tracer(a.indices ∪ b.indices)
```

```
/(a :: Tracer, b :: Real) = Tracer(a.indices)
```

```
sign(a :: Tracer) = Tracer() # zero derivatives
```

# Pocket pattern detection

```
julia> f(x) = [x[1] * x[2] * sign(x[3]), sign(x[3]) * x[4] / 2];
```

```
julia> x = Tracer.(Set.([1, 2, 3, 4]))
```

```
4-element Vector{Tracer}:
```

```
Tracer(Set([1]))
```

```
Tracer(Set([2]))
```

```
Tracer(Set([3]))
```

```
Tracer(Set([4]))
```

```
julia> f(x)
```

```
2-element Vector{Tracer}:
```

```
Tracer(Set([2, 1]))
```

```
Tracer(Set([4]))
```

# Coloring for Jacobians

## Matrix problem

Orthogonal partition of the columns of  $A$ .

If  $A_{ij_1} \neq 0$  and  $A_{i,j_2} \neq 0$ , then columns  $j_1$  and  $j_2$  are in different groups  $c(j_1) \neq c(j_2)$

These are equivalent<sup>4</sup> if we define the graph representation

$$\mathcal{E} = \{(i, j) \in \mathcal{I} \times \mathcal{J} : A_{i,j} \neq 0\}$$

## Graph problem

Partial distance-2 coloring of a bipartite graph  $\mathcal{G} = (\mathcal{I} \cup \mathcal{J}, \mathcal{E})$

If  $(i, j_1) \in \mathcal{E}$  and  $(i, j_2) \in \mathcal{E}$ , then vertices  $j_1$  and  $j_2$  have different colors  $c(j_1) \neq c(j_2)$ .

---

<sup>4</sup>Gebremedhin et al. (2005)

# Coloring for Jacobians, illustrated

Coloring of intersection graph / distance-2 coloring of bipartite graph

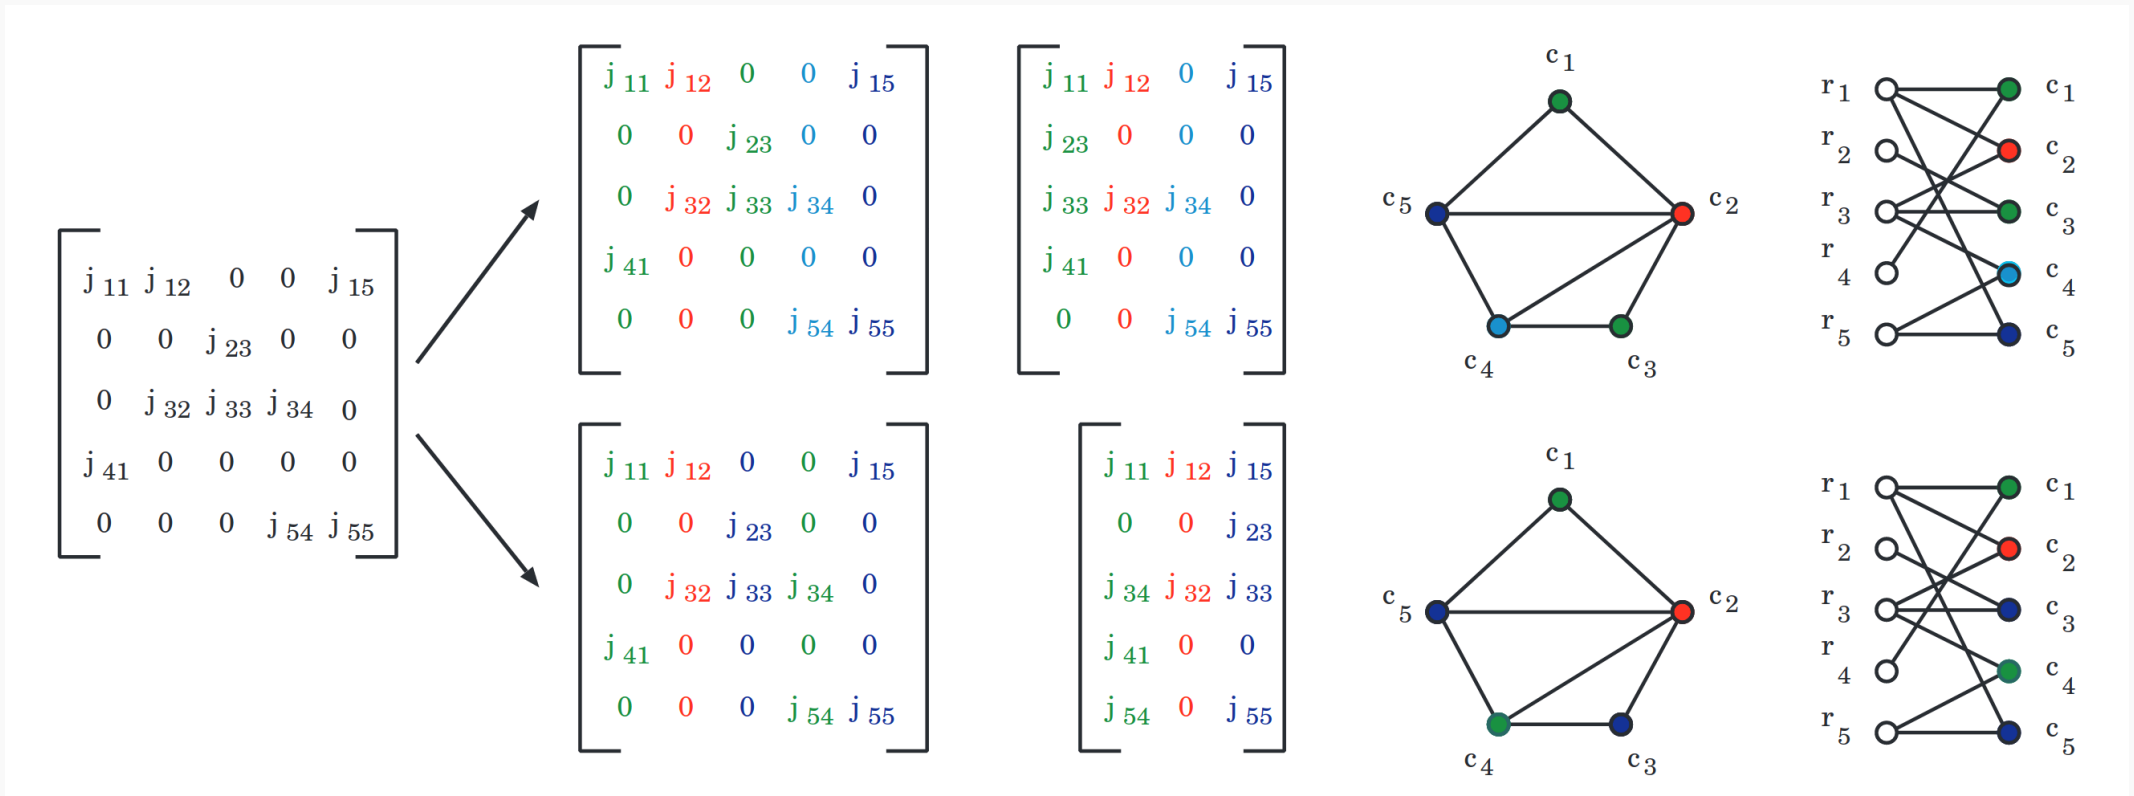


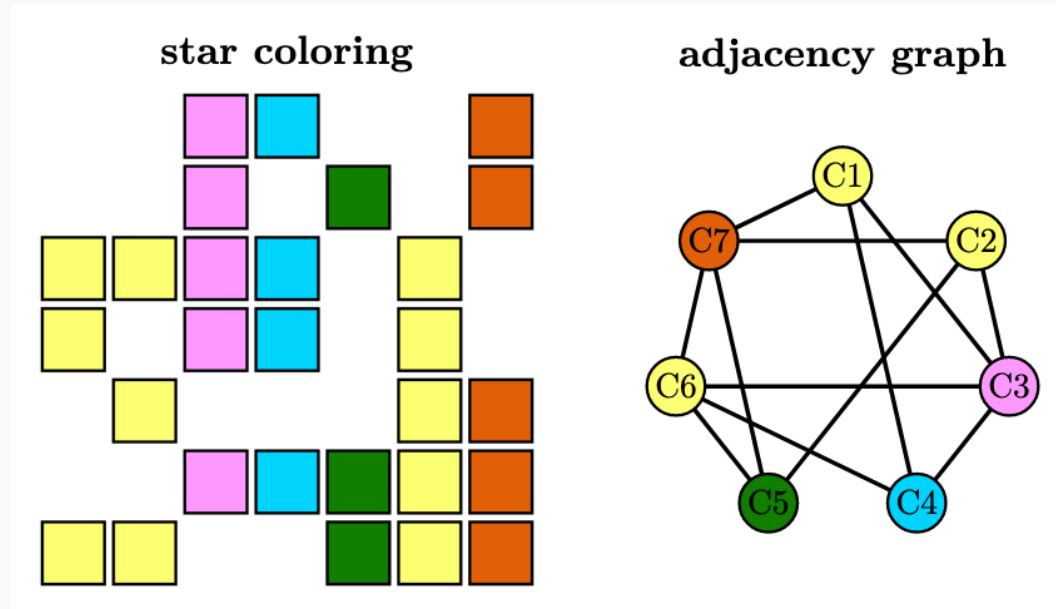
Figure 4: Gebremedhin et al. (2005)



# Coloring for Hessians

What if our matrix has structure, like  $A_{i,j} = A_{j,i}$ ?

We can compute a slightly different coloring<sup>5</sup> with fewer colors.



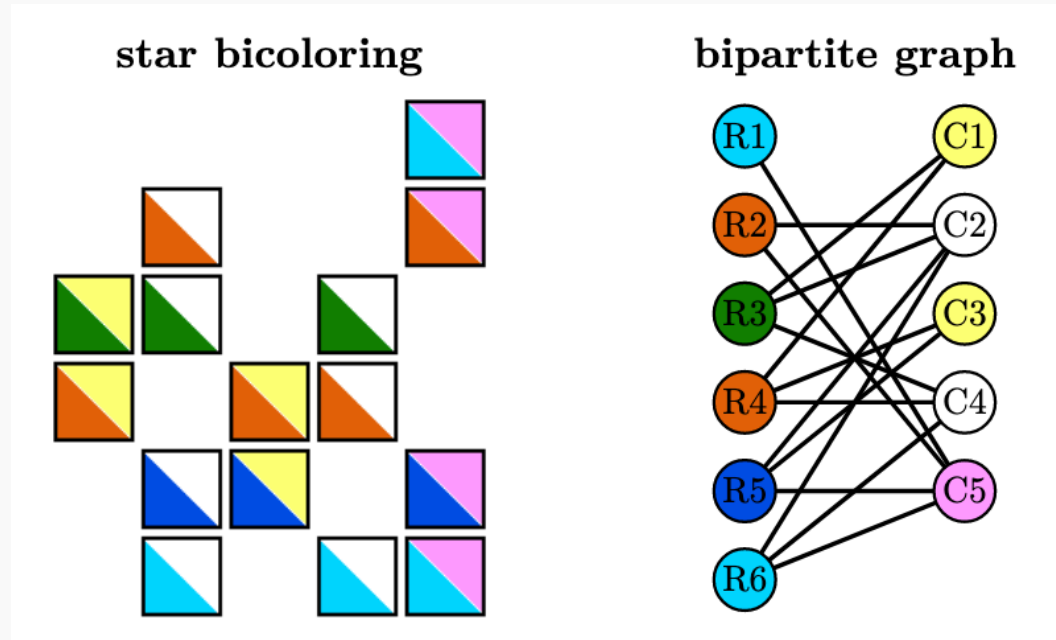
---

<sup>5</sup>Coleman & Moré (1984)

# Coloring for bidirectional Jacobians

What if the columns are not orthogonal enough?

We can use both rows and columns<sup>6</sup> inside our coloring.

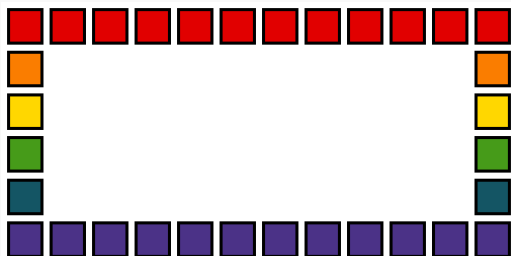


---



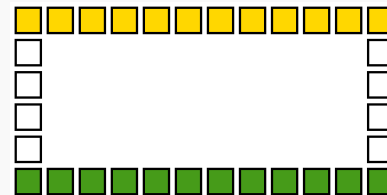
<sup>6</sup>Coleman & Verma (1998)

# Benefits of bidirectional coloring

Compute Jacobians with a dense row **and** a dense column, using forward-mode AD + reverse-mode AD.



Unidirectional



Bidirectional

# From bidirectional to symmetric [new]

To color the rows and columns of  $J$ , color the columns of  $H = \begin{pmatrix} 0 & J \\ J & 0 \end{pmatrix}$

It sounds simple, but:

- Some colors may be redundant
- Detecting these is tightly linked to the two-colored structures
- Efficient decompression requires lots of preprocessing

Explanations and benchmarks in Montoison et al. (2025)

# The sharp bits

## Pattern detection

- Local versus global sparsity
- Control flow
- Linear and nonlinear interactions



## Coloring

- Only heuristic algorithms
- Vertex ordering matters a lot

# Implementation

---

# AD in Python & Julia



# Interfaces for experimenting [new]



Figure 8: In Python, Keras supports Tensorflow, PyTorch and JAX.




Figure 9: In Julia, 14 AD backends inside `DifferentiationInterface.jl`

Once we have a common syntax, we can do more!




# One API to rule them all


Decouple the scientific libraries from the AD underneath.



(a) Python



(b) Julia (before DI)



(c) Julia (now)

# Previous implementations of sparse AD

- In low-level programming languages (C, Fortran)
- In closed-source languages (Matlab)
- In domain-specific languages (AMPL, CasADi)

Basically nothing in Python (either in JAX or PyTorch).

First drafts in Julia for scientific machine learning, but severely limited: single-backend, slow.

# A modern sparse AD ecosystem [new]

Independent packages working together:

- **Step 1:** `SparseConnectivityTracer.jl` (Hill & Dalle, 2025)
- **Steps 2 & 4:** `SparseMatrixColorings.jl` (Montoison et al., 2025)
- **Step 3:** `Differentiationinterface.jl` (Dalle & Hill, 2025)

	SCT.jl	SMC.jl	DI.jl
lines of code	5202	5184	19980
indirect dependents	461	487	896
downloads / month	7.8k	9.7k	33k

Compatible with generic code!

# Impact

Users already include...

- Scientific computing: `SciML` (Julia's `scipy`)
  - Differential equations
  - Nonlinear solvers
  - Optimization
- Probabilistic programming: `Turing.jl`
- Symbolic regression: `PySR`

This is the part where things go sideways.

# Perspectives

- GPU-compatible pattern detection and coloring
- Pattern detection in JAX with program transformations
- New, unsuspected applications “just because we can”

# Going further

On general AD:

- Baydin et al. (2018)
- Margossian (2019)
- Blondel & Roulet (2024)

On sparse AD:

- Gebremedhin et al. (2005)
- Griewank & Walther (2008)
- Hill et al. (2025)

# Bibliography

- Baydin, A. G., Pearlmutter, B. A., Radul, A. A., & Siskind, J. M. (2018). Automatic Differentiation in Machine Learning: A Survey. *Journal of Machine Learning Research*, 18(153), 1–43. <http://jmlr.org/papers/v18/17-468.html>
- Blondel, M., & Roulet, V. (2024, July). *The Elements of Differentiable Programming*. arXiv. <https://doi.org/10.48550/arXiv.2403.14606>
- Coleman, T. F., & Moré, J. J. (1984). Estimation of Sparse Hessian Matrices and Graph Coloring Problems. *Mathematical Programming*, 28(3), 243–270. <https://doi.org/10.1007/BF02612334>
- Coleman, T. F., & Verma, A. (1998). The Efficient Computation of Sparse Jacobian Matrices Using Automatic Differentiation. *SIAM Journal on Scientific Computing*, 19(4), 1210–1233. <https://doi.org/10.1137/S1064827595295349>
- Curtis, A. R., Powell, M. J. D., & Reid, J. K. (1974). On the Estimation of Sparse Jacobian Matrices. *IMA Journal of Applied Mathematics*, 13(1), 117–119. <https://doi.org/10.1093/imamat/13.1.117>
- Dalle, G., & Hill, A. (2025, May). *A Common Interface for Automatic Differentiation*. arXiv. <https://doi.org/10.48550/arXiv.2505.05542>
- Gebremedhin, A. H., Manne, F., & Pothen, A. (2005). What Color Is Your Jacobian? Graph Coloring for Computing Derivatives. *SIAM Review*, 47(4), 629–705. <https://doi.org/10/cmws4>
- Griewank, A., & Walther, A. (2008). *Evaluating Derivatives: Principles and Techniques of Algorithmic Differentiation* (2nd ed). Society for Industrial and Applied Mathematics. <https://epubs.siam.org/doi/book/10.1137/1.9780898717761>
- Hill, A., & Dalle, G. (2025). Sparser, Better, Faster, Stronger: Sparsity Detection for Efficient Automatic Differentiation. *Transactions on Machine Learning Research*. <https://openreview.net/forum?id=GtXSN52nIW>
- Hill, A., Dalle, G., & Montoison, A. (2025, April). An Illustrated Guide to Automatic Sparse Differentiation. *ICLR Blogposts 2025*. <https://iclr-blogposts.github.io/2025/blog/sparse-autodiff/>



- Margossian, C. C. (2019). A Review of Automatic Differentiation and Its Efficient Implementation. *Wires Data Mining and Knowledge Discovery*, 9(4), e1305. <https://doi.org/10.1002/widm.1305>
- Montoisson, A., Dalle, G., & Gebremedhin, A. (2025, May). *Revisiting Sparse Matrix Coloring and Bicoloring*. arXiv. <https://doi.org/10.48550/arXiv.2505.07308>
- Pearlmutter, B. A. (1994). Fast Exact Multiplication by the Hessian. *Neural Computation*, 6(1), 147–160. <https://doi.org/10.1162/neco.1994.6.1.147>